



SAPIENZA  
UNIVERSITÀ DI ROMA

Faculty of Information Engineering, Informatics and  
Statistics  
Department of Computer Science

# Distributed Systems

**Author:**  
Simone Lidonnici

24 september 2025

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
1.1	Defining a distributed system . . . . .	1
1.2	Computation . . . . .	1
1.3	Monitoring computations . . . . .	3

# 1

## Introduction

### 1.1 Defining a distributed system

We define a **distributed system** as a collection of processes  $p_1, p_2, \dots, p_n$  that run on different computers and cooperate to solve a problem. The processes communicate using **channels** and we assume the system is fully connected, meaning that every pair of processes can exchange messages between them. The channels are reliable, in the sense that the message arrives, but may be delivered out of order.

The simple model for a distributed system is called **asynchronous**, that have no upper bound on the speed of processes and no upper bound for the delay of a message. If these upper bounds exist the system is called **synchronous**. The second one is stronger, in the sense that we do more assumptions and this means that every program that runs on a synchronous system can run on an asynchronous system (the opposite can be possible but not sure).

A distributed system can have different properties:

- **Consistency**: every part of the system has the same information at every time
- **Availability**: the information is available at every time
- **Partition tolerance**: if one part of the system goes offline the system can continue to run

Every type of distributed system can have up to 2 of these properties.

### 1.2 Computation

We describe the execution of a program on a distributed system as a collection of processes. Every process is defined as a sequence of **events**. The events can be internal or involve communication like the events **send(m)** and **receive(m)**.

We label an event with the notation:

$$e_i^k$$

in which  $i$  represents the index of the process  $p_i$  and  $k$  the order of the event for that specific process.

### Local and global history

The **local history** of a process  $p_i$  is a sequence of events  $h_i = e_i^1 e_i^2 \dots$  that represent the sequential execution of events in the process. We use  $h_i^k$  to represent the sequence of the first  $k$  events in the process  $p_i$ .

The **global history** of the computation is a set that contains all events:

$$H = h_1 \cup h_2 \cup \dots \cup h_n$$

The global history gives us no information about the time in which these events are executed, because in an asynchronous system there is no global clock.


### Relation of cause-effect

Events can be labeled based of the notion of **cause-effect**, defining a relation (with symbol  $\rightarrow$ ) between two events such that:

- $\forall e_i^k, e_i^l \in h_i \wedge k < l \implies e_i^k \rightarrow e_i^l$
- $e_i = \text{send}(\mathbf{m}) \wedge e_j = \text{receive}(\mathbf{m}) \implies e_i \rightarrow e_j$
- $e \rightarrow e' \wedge e' \rightarrow e'' \implies e \rightarrow e''$  (Transitive)

This relation mean that  $e \rightarrow e'$  if  $e$  causally precedes  $e'$ , so the computation of  $e'$  is influenced by  $e$ . Two events can be unrelated, so neither  $e \rightarrow e'$  nor  $e' \rightarrow e$ . We call this pair of events as **concurrent** and write them as  $e || e'$ .

We can graphically represent a computation with a space-time diagram like this:



An arrow from  $p_1$  to  $p_2$  means that  $p_1$  sends a message to  $p_2$ .

### Run

A **run** is a total order of all events in the global history, consistent with each local history, so the events in history  $h_i$  appear in the same order in  $R$ :

$$R = e_1^1 e_2^1 \dots$$

A single program can have many different runs because some events are unrelated.

## 1.3 Monitoring computations

### Local and global state

We denote  $\sigma_i^k$  the **local state** of a process  $p_i$  after the event  $e_i^k$ .  
The **global state** of the computation is an  $n$ -tuple of local states:

$$\Sigma = (\sigma_1, \sigma_2, \dots, \sigma_n)$$

### Cut

A **cut** is a collection of local histories:


$$C = \langle h_1^{k_1}, h_2^{k_2}, \dots, h_n^{k_n} \rangle$$

To monitor the computation or to compute a global problem (for example knowing if the system is in deadlock) we add a process  $p_0$ .

The first idea is to make  $p_0$  send a message to every process to which a process  $p_i$  will respond with the current state  $\sigma_i$ . After all the response  $p_0$  can construct a global state, that defines a cut.

#### Example:


Considering the follows space-time diagram with the cut  $C$  (the dashed arrow represent when the process sends the response to  $p_0$ ):



If we use the responses to create a graph, based on the states, we can say that:

- $p_1$  is going to send a response to  $p_2$
- $p_2$  is going to send a response to  $p_3$
- $p_3$  is going to send a response to  $p_1$

So the graph generated by  $p_0$  will be:



It seems the system has a deadlock, but if we watch carefully we can see that this is not true. The problem is that the global state defined by the cut  $C$  could never happen during a computation.

### Consistent cut

A cut is **consistent** if only if:

$$\forall e \rightarrow e' \wedge e' \in C \implies e \in C$$

So a cut is consistent if the global state generated by the cut could be possible during a computation.

If we use  $p_0$  only to receive messages and we make every process notify the events to  $p_0$  with a timestamp associated,  $p_0$  can reorder the events to create a run. This is true only if there is a global clock, which is not true. To overcome this problem we have to create a local clock for every process and update it in a consistent way.

### Clock condition

A run is consistent if only if follows the **clock condition**:

$$\forall e \rightarrow e' \implies TS(e) < TS(e')$$

If  $TS(e) < TS(e')$  is possible but not guaranteed that  $e \rightarrow e'$ .

### Local clock

We define the **local clock** of a process as follows:

$$LC(e_i) = \begin{cases} LC + 1 & e_i \text{ is internal or send} \\ \max(LC, TS(m)) + 1 & e_i \text{ is a receive(m)} \end{cases}$$